


CSC4200/5200 – COMPUTER NETWORKING

RELIABLE DELIVERY – PART 1

Instructor: Susmit Shannigrahi
sshannigrahi@tnitech.edu



Frames – bag of bits



- Sending side – encapsulation, add error check bits, flow control
- Receiving side – extract frames, check for error, flow control

Reliable Delivery

- Frames might get lost
 - Too many bits lost
 - Clock did not sync properly
 - Error detected but the report got lost
- Can we build links that does not have errors?
 - Not possible
- How about all those error correction stuff we learned?
 - Can we add them to frames?
 - We could, but think of the overhead
 - What happens when the entire frame is lost?

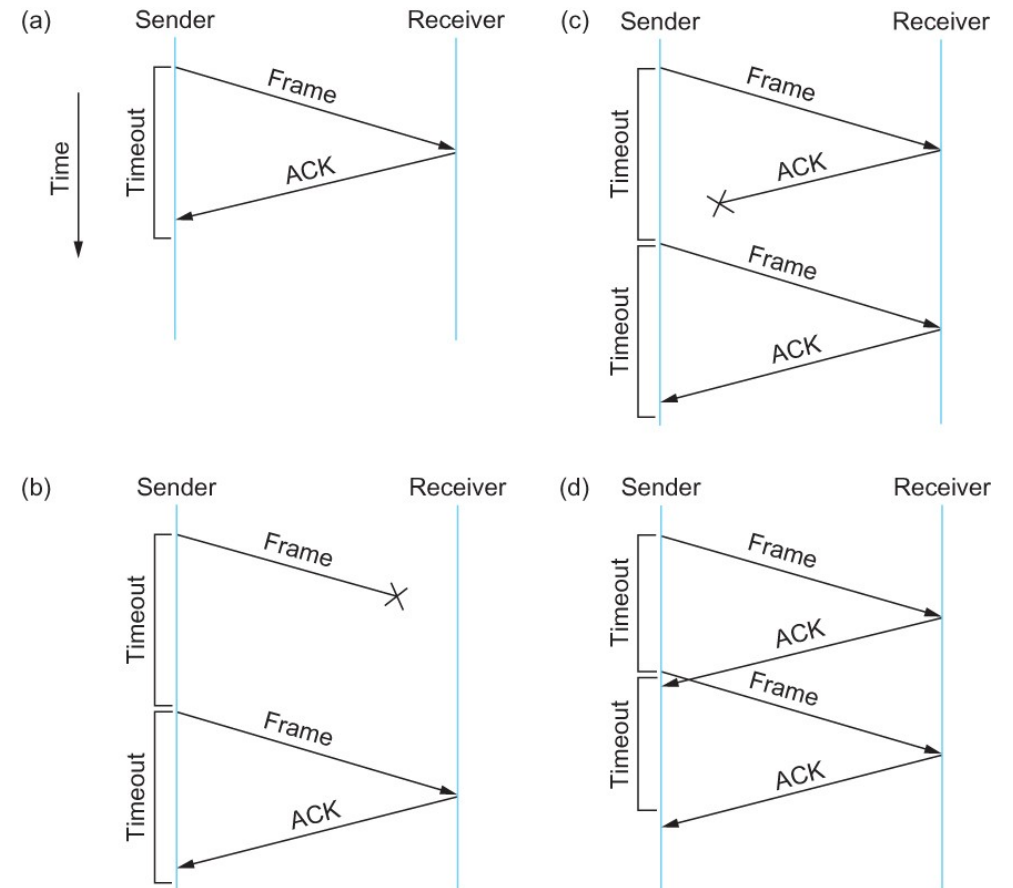
Frames – bag of bits



- Sending side – encapsulation, add error check bits, flow control
- Receiving side – extract frames, check for error, flow control


Stop and Wait

- Sender sends a frame, sets a timeout (e.g., 1 sec)
- Receiver receives the frame, sends an ACK
- Sender
 - sends the next frame on ACK
 - retransmits the same frame if timeout happens
- **Spot the bugs in the protocol**




Stop and Wait – Bugs (C and D)


(a)




(c)



(b)




(d)

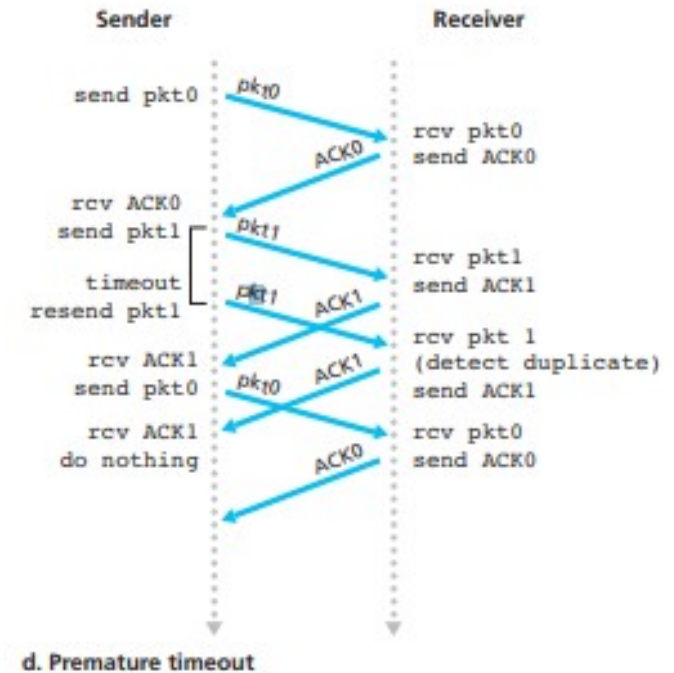
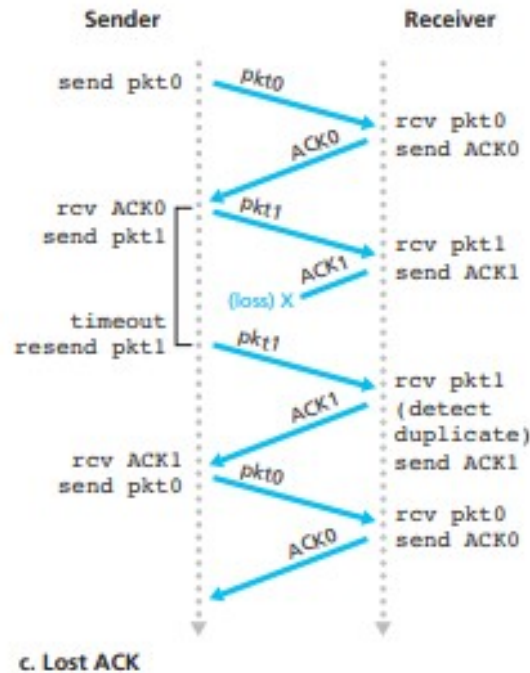




Stop and Wait – How to fix the bug?

Hint: Uniquely identify each packet



Stop and Wait v2




Stop and Wait - V2 Problems

- Sender sets a timeout to wait for an ACK
 - Too small – retransmissions
 - Too large – long wait if frames are lost


- Solution:

- Keep a running average of Round Trip Time
 - $\text{EstimatedRTT} = (1 - \alpha) \cdot \text{EstimatedRTT} + \alpha \cdot \text{SampleRTT}$
 - $\text{Timeout} = 2 \cdot \text{EstimatedRTT}$
 - Value of $\alpha = 0.125$
- Where does α come from? RFC 6928 (for now)




Stop and Wait – How to fix the bug?

Hint: Uniquely identify each packet



Stop and Wait – How does it perform?

- Bandwidth (R) = 1Gbps
- Packet size (L) = 1000 bytes
- RTT = 30ms
- $T_{\text{trans}} = L/R = 8000\text{bits}/10^9\text{bits/sec} = 8\text{microsecond}$
- $T_{\text{prop}} = 15\text{ms}$
- Total Delay = 15.008 ms




a. A stop-and-wait protocol in operation

Kurose/Ross

Stop and Wait – How does it perform?


- Sender transmits for only 0.008 ms in 30.008ms
- Utilization = $0.008/30.008 = 0.00027$
- One bit at a time
- Worse when loss happens!



Kurose/Ross

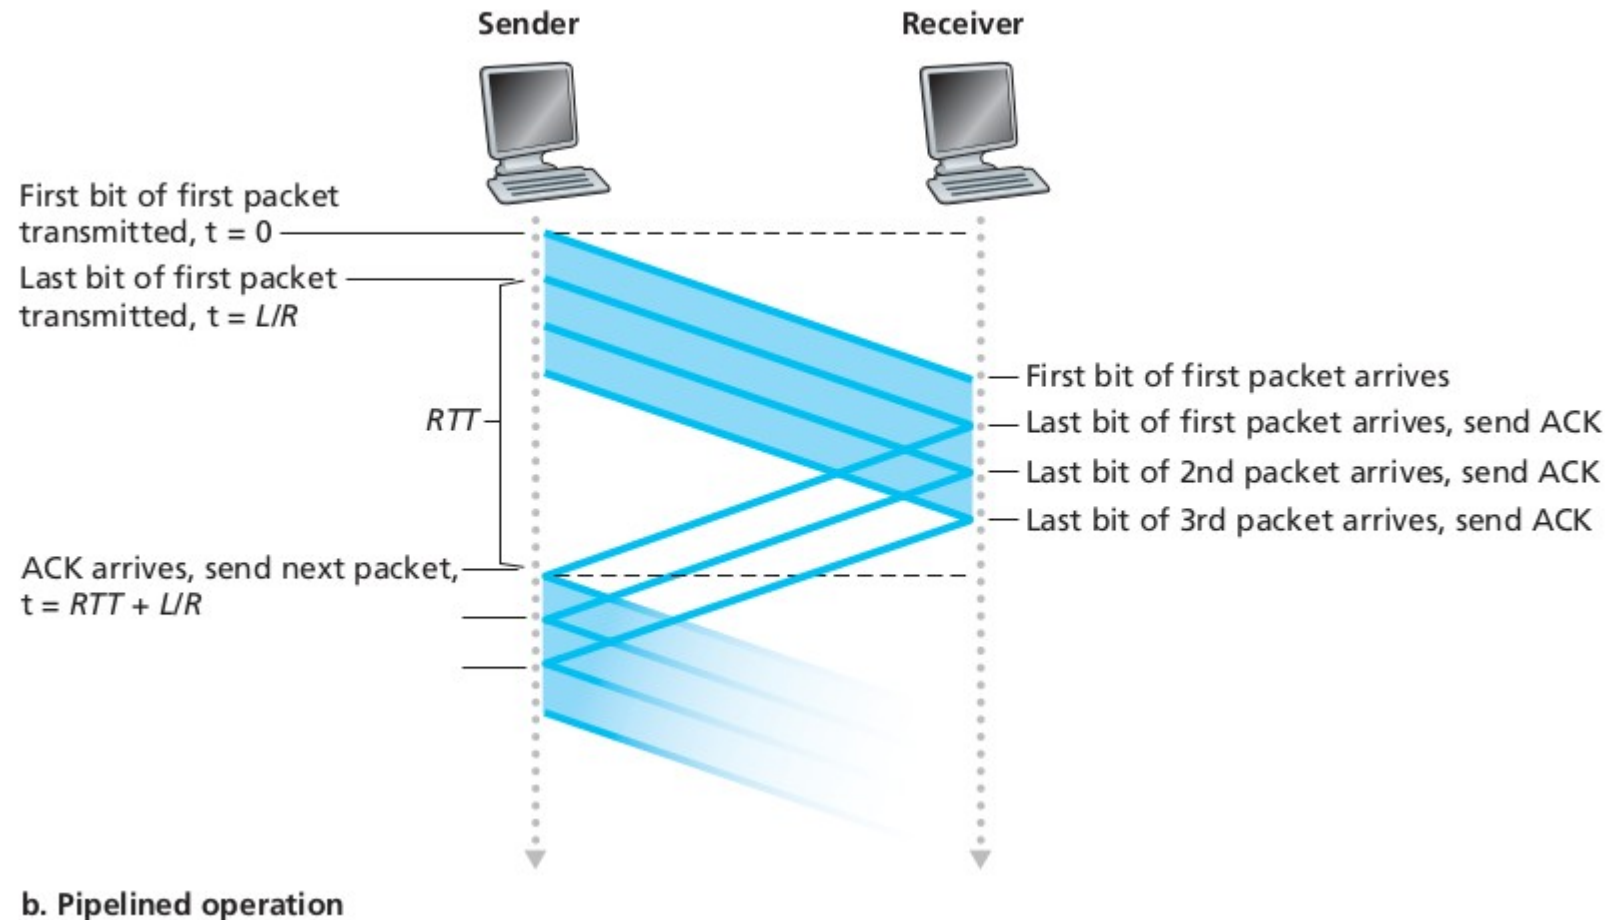
Sliding window to the rescue!

Utilization = $0.008 * 3 / 30.008 = 0.00079$ (3 times increase)



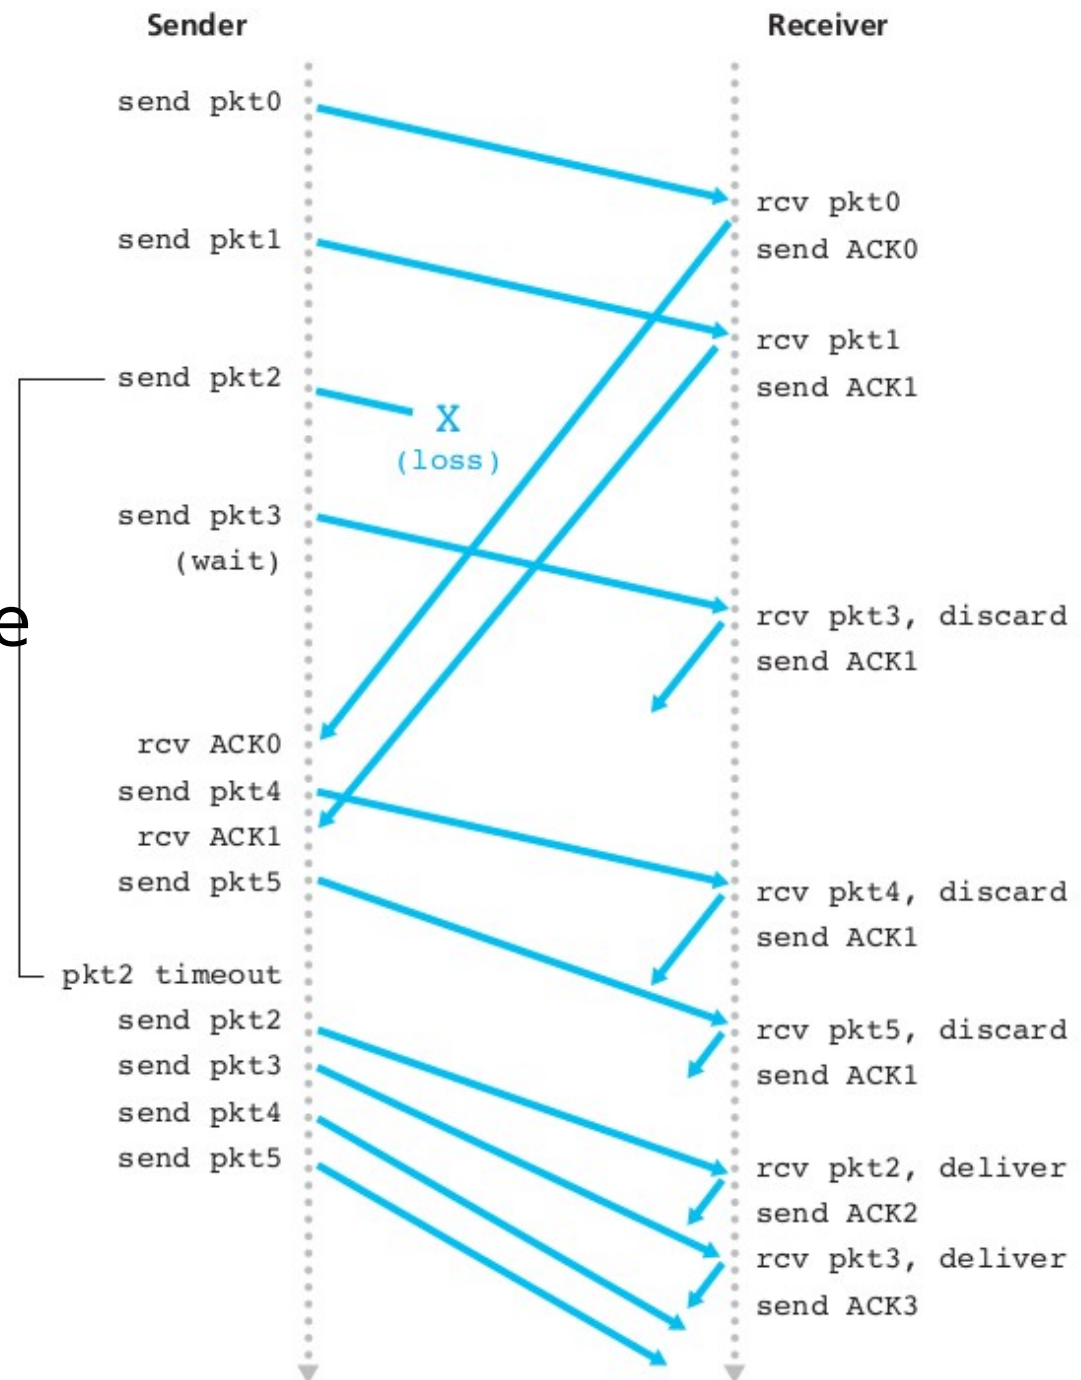
Sliding window to the rescue!

Utilization = $0.008 * 3 / 30.008 = 0.00079$ (3 times increase)



Go-Back-N

- See the problem?
- Can not move forward until all previous packets are acknowledged




Selective Repeat


http://www.exa.unicen.edu.ar/catedras/comdat1/material/Filminas3_Practico3.swf

- Receiver:
 - Individually acks all packets
 - Buffers packets as necessary
 - Buffer packets until lost packets are received
- Sender:
 - Resend packets (**only**) for which ACK not received
 - Timer for each unACKed packet
 - Can send only n packets

Sliding window - Selective Repeat



a. Sender view of sequence numbers



b. Receiver view of sequence numbers

Key:

Already ACK'd	Usable, not yet sent
Sent, not yet ACK'd	Not usable

Key:

Out of order (buffered) but already ACK'd	Acceptable (within window)
Expected, not yet received	Not usable

Sender

Receiver

pkt0 sent

0 1 2 3 4 5 6 7 8 9

pkt1 sent

0 1 2 3 4 5 6 7 8 9

pkt2 sent

0 1 2 3 4 5 6 7 8 9

pkt3 sent, window full

0 1 2 3 4 5 6 7 8 9

ACK0 rcvd, pkt4 sent

0 1 2 3 4 5 6 7 8 9

ACK1 rcvd, pkt5 sent

0 1 2 3 4 5 6 7 8 9

pkt2 TIMEOUT, pkt2
resent

0 1 2 3 4 5 6 7 8 9

ACK3 rcvd, nothing sent

0 1 2 3 4 5 6 7 8 9

X
(loss)

pkt0 rcvd, delivered, ACK0 sent

0 1 2 3 4 5 6 7 8 9

pkt1 rcvd, delivered, ACK1 sent

0 1 2 3 4 5 6 7 8 9

pkt3 rcvd, buffered, ACK3 sent

0 1 2 3 4 5 6 7 8 9

pkt4 rcvd, buffered, ACK4 sent

0 1 2 3 4 5 6 7 8 9

pkt5 rcvd; buffered, ACK5 sent

0 1 2 3 4 5 6 7 8 9

pkt2 rcvd, pkt2, pkt3, pkt4, pkt5
delivered, ACK2 sent

0 1 2 3 4 5 6 7 8 9

Sliding window - Selective Repeat - LOSS

- **Sender:**

- Data received, if next to-be-sent-packet's seq # within window, send. Else, buffer or return to application.
- Timeout: Each packet has its own timer. resend the packet
- ACK received: Mark received, Advance window to next unacked seq # if ack for send_base

- **Receiver, packet (n)**

- Sequence between recv_base, recv_base + N - 1, send ack (n)
- Out of order: buffer
- In-order or closes gap – deliver to application
- Packet within $<\text{recv_base}-N, \text{recv_base}-1>$, ACK(n)
- Otherwise: Ignore

Issues with Sliding Window Protocol

- When timeout occurs, the amount of data in transit decreases
 - Since the sender is unable to advance its window
- When the packet loss occurs, this scheme is no longer keeping the pipe full
 - The longer it takes to notice that a packet loss has occurred, the more severe the problem becomes
- How to improve this
 - Negative Acknowledgement (NAK)
 - Additional Acknowledgement
 - Selective Acknowledgement (SAK)

Next Steps

- Read Through - Chapter 2.5.2
- Ethernet and WiFi
- Project 1 due on the 20th
- Project 2 and homework 2 will post on the 20th