


CSC4200/5200 – COMPUTER NETWORKING

Instructor: Susmit Shannigrahi

BGP - CONTINUED

sshannigrahi@tnitech.edu


GTA: dereddick42@students.tnitech.edu



So far...

- Routing
How do we scale routing?

Internet now



Hierarchical routing - Policy

scale: with 600 million destinations:

- can't store all dest's in routing tables!
- routing table exchange would swamp links!

administrative autonomy


- internet = network of networks
- each network admin may want to control routing in its own network

Autonomous systems (ASes)

- AS
 - A set of routers under a single technical administration
 - Uses IGP within the AS to route packets
 - Uses BGP between Ases to route packets
- What happens inside an AS stays within that AS!
 - That is, AS decides routing metrics internally


Status of ASNs

Status	AS Pool	16-bit	32-bit
IETF Reserved	95033874	1042	95032832
IANA Unallocated Pool	4199828976	0	4199828976
Allocated	104446	64494	39952
RIR Data			
AFRINIC	2302	1278	1024
APNIC	19093	8539	10554
ARIN	31567	25522	6045
RIPE NCC	39453	25729	13724
LACNIC	12031	3426	8605



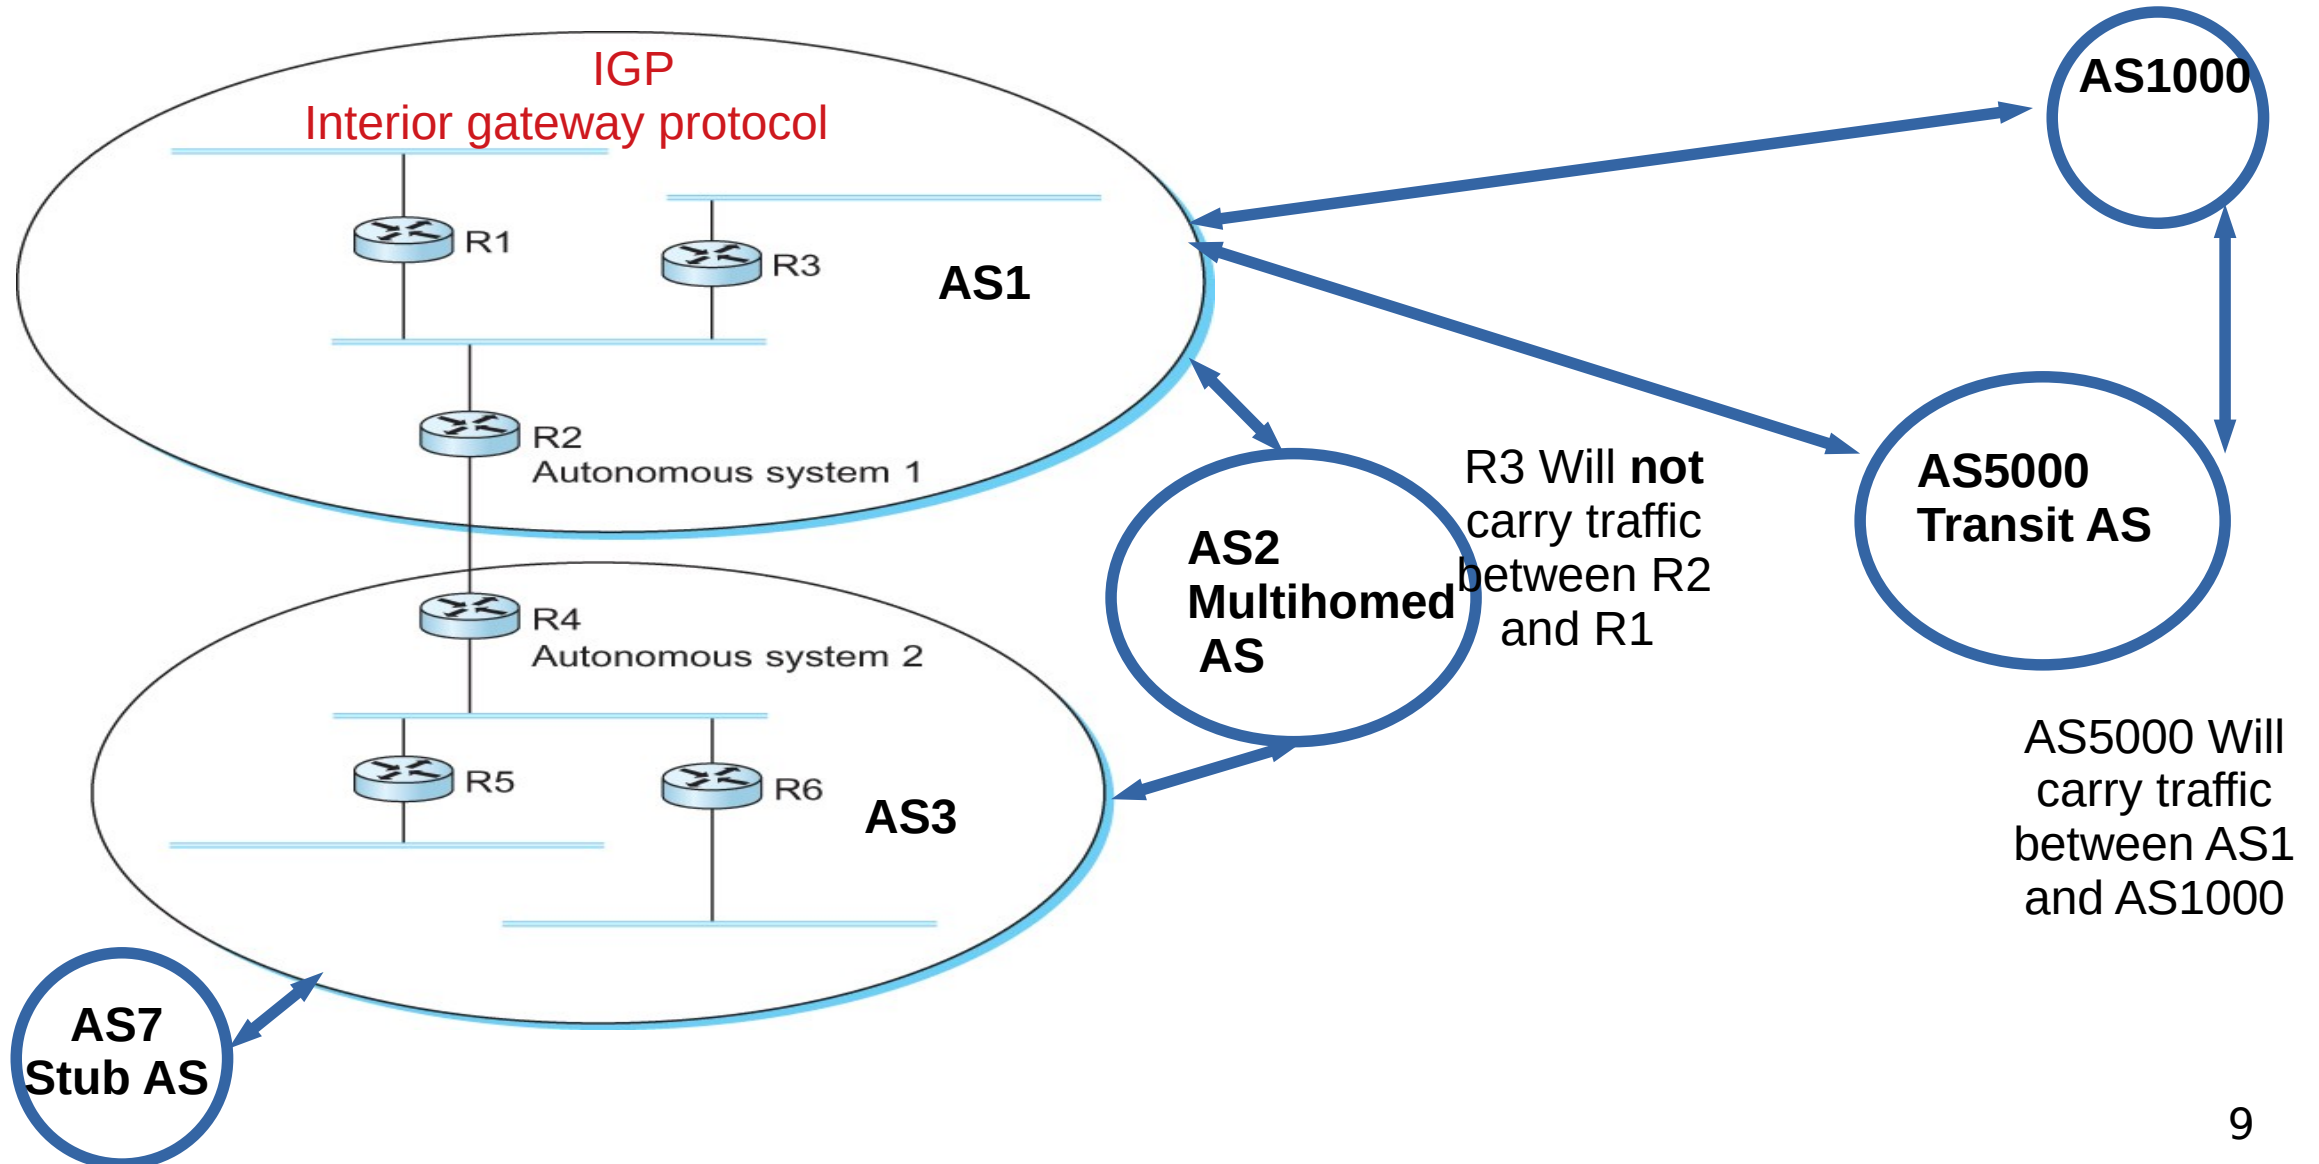
<http://www.potaroo.net/tools/asn32/>

Interdomain Routing



A network with four autonomous systems

BGP-4: Border Gateway Protocol



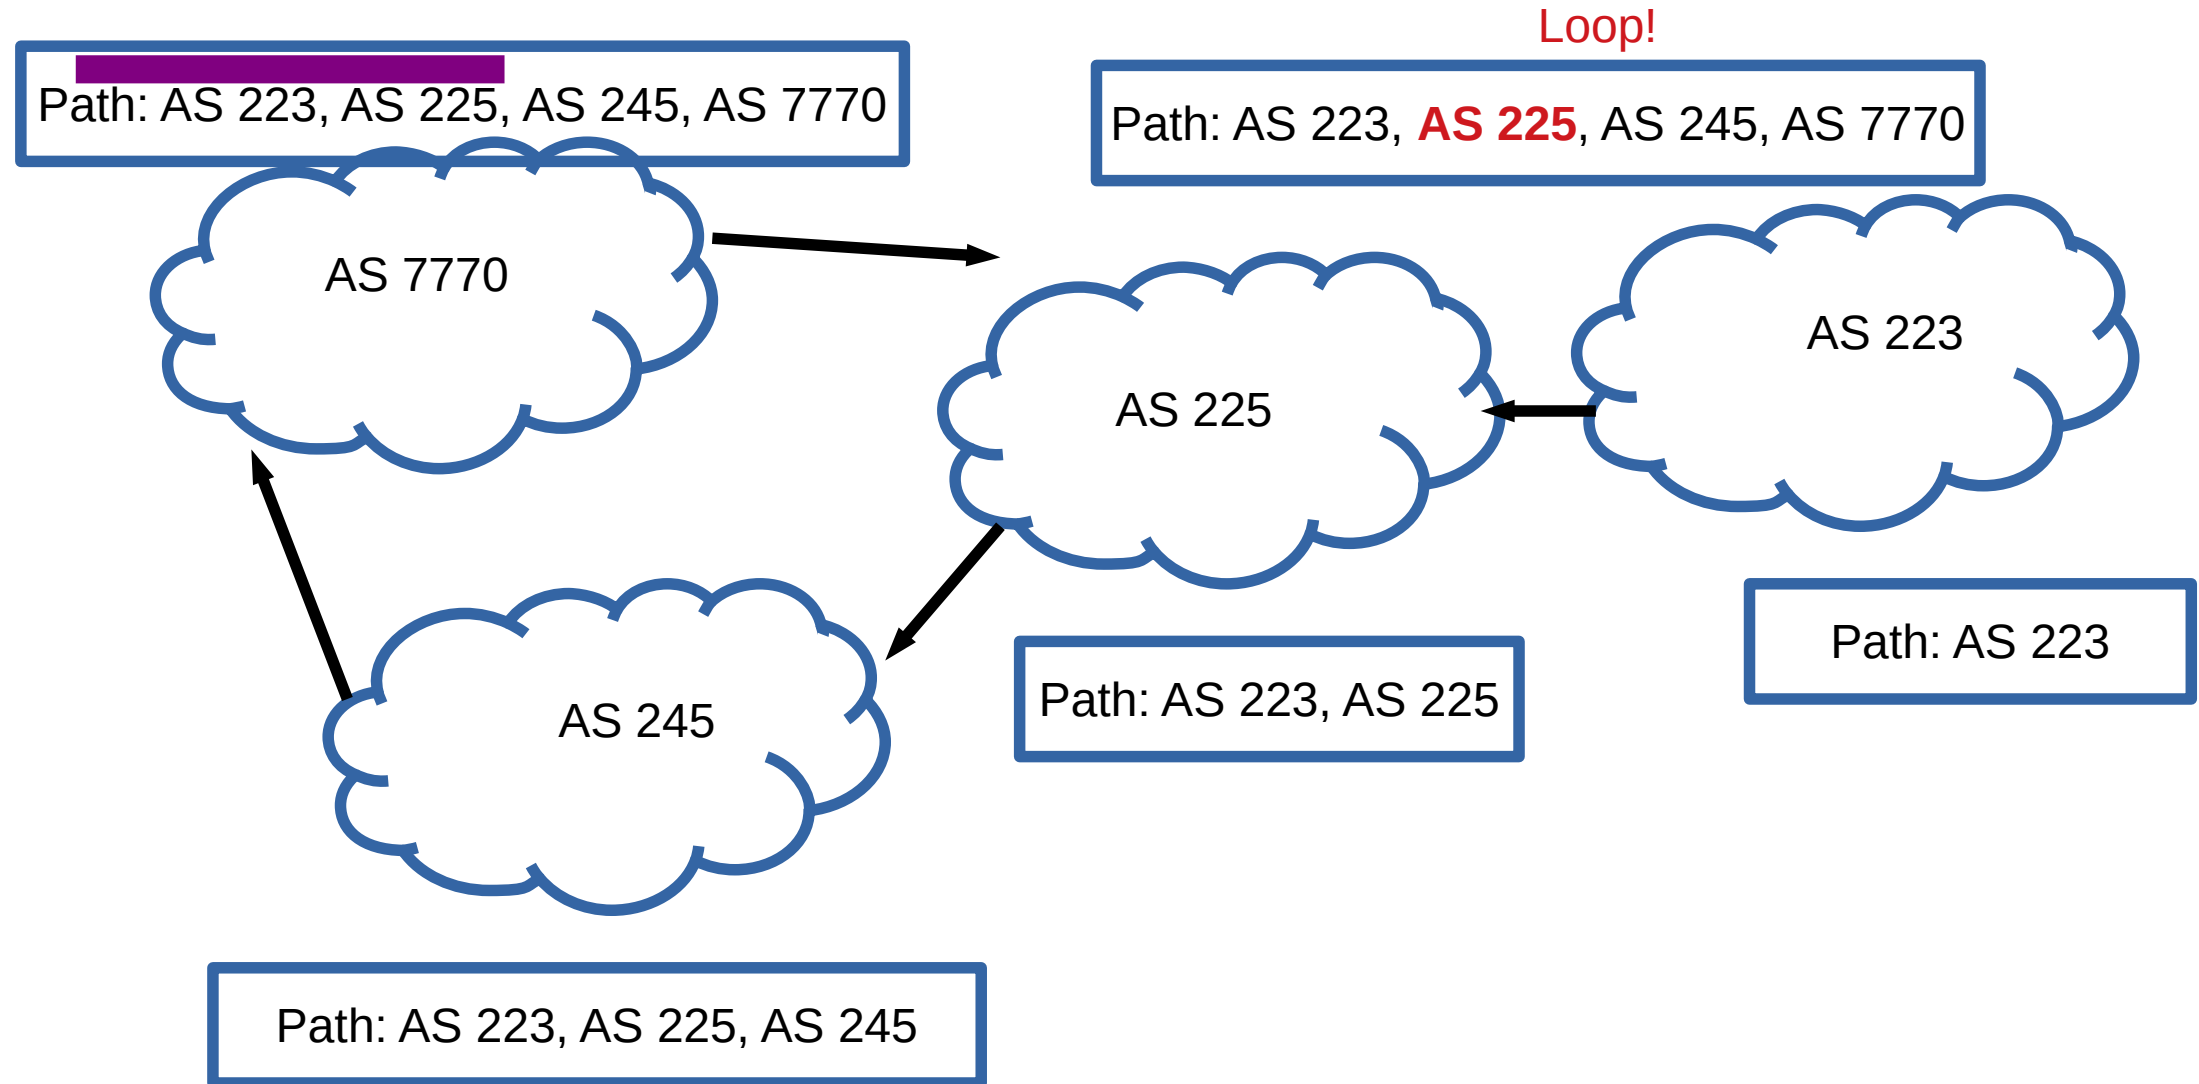
BGP - goals

- The goal of Inter-domain routing is to find **any path** to the intended destination that is **loop free**
 - **We are concerned with reachability than optimality**
 - Finding path anywhere close to optimal is considered to be a great achievement
- Why?

BGP: Path vector protocol

- Send the whole path with the routing update
- Loops are detected if an AS finds itself in the path
 - Reject if so
 - Accept otherwise
- Add self to the path and advertise to the neighbors
- Advantage: No loops, Local decision before advertising

BGP: Path vector protocol



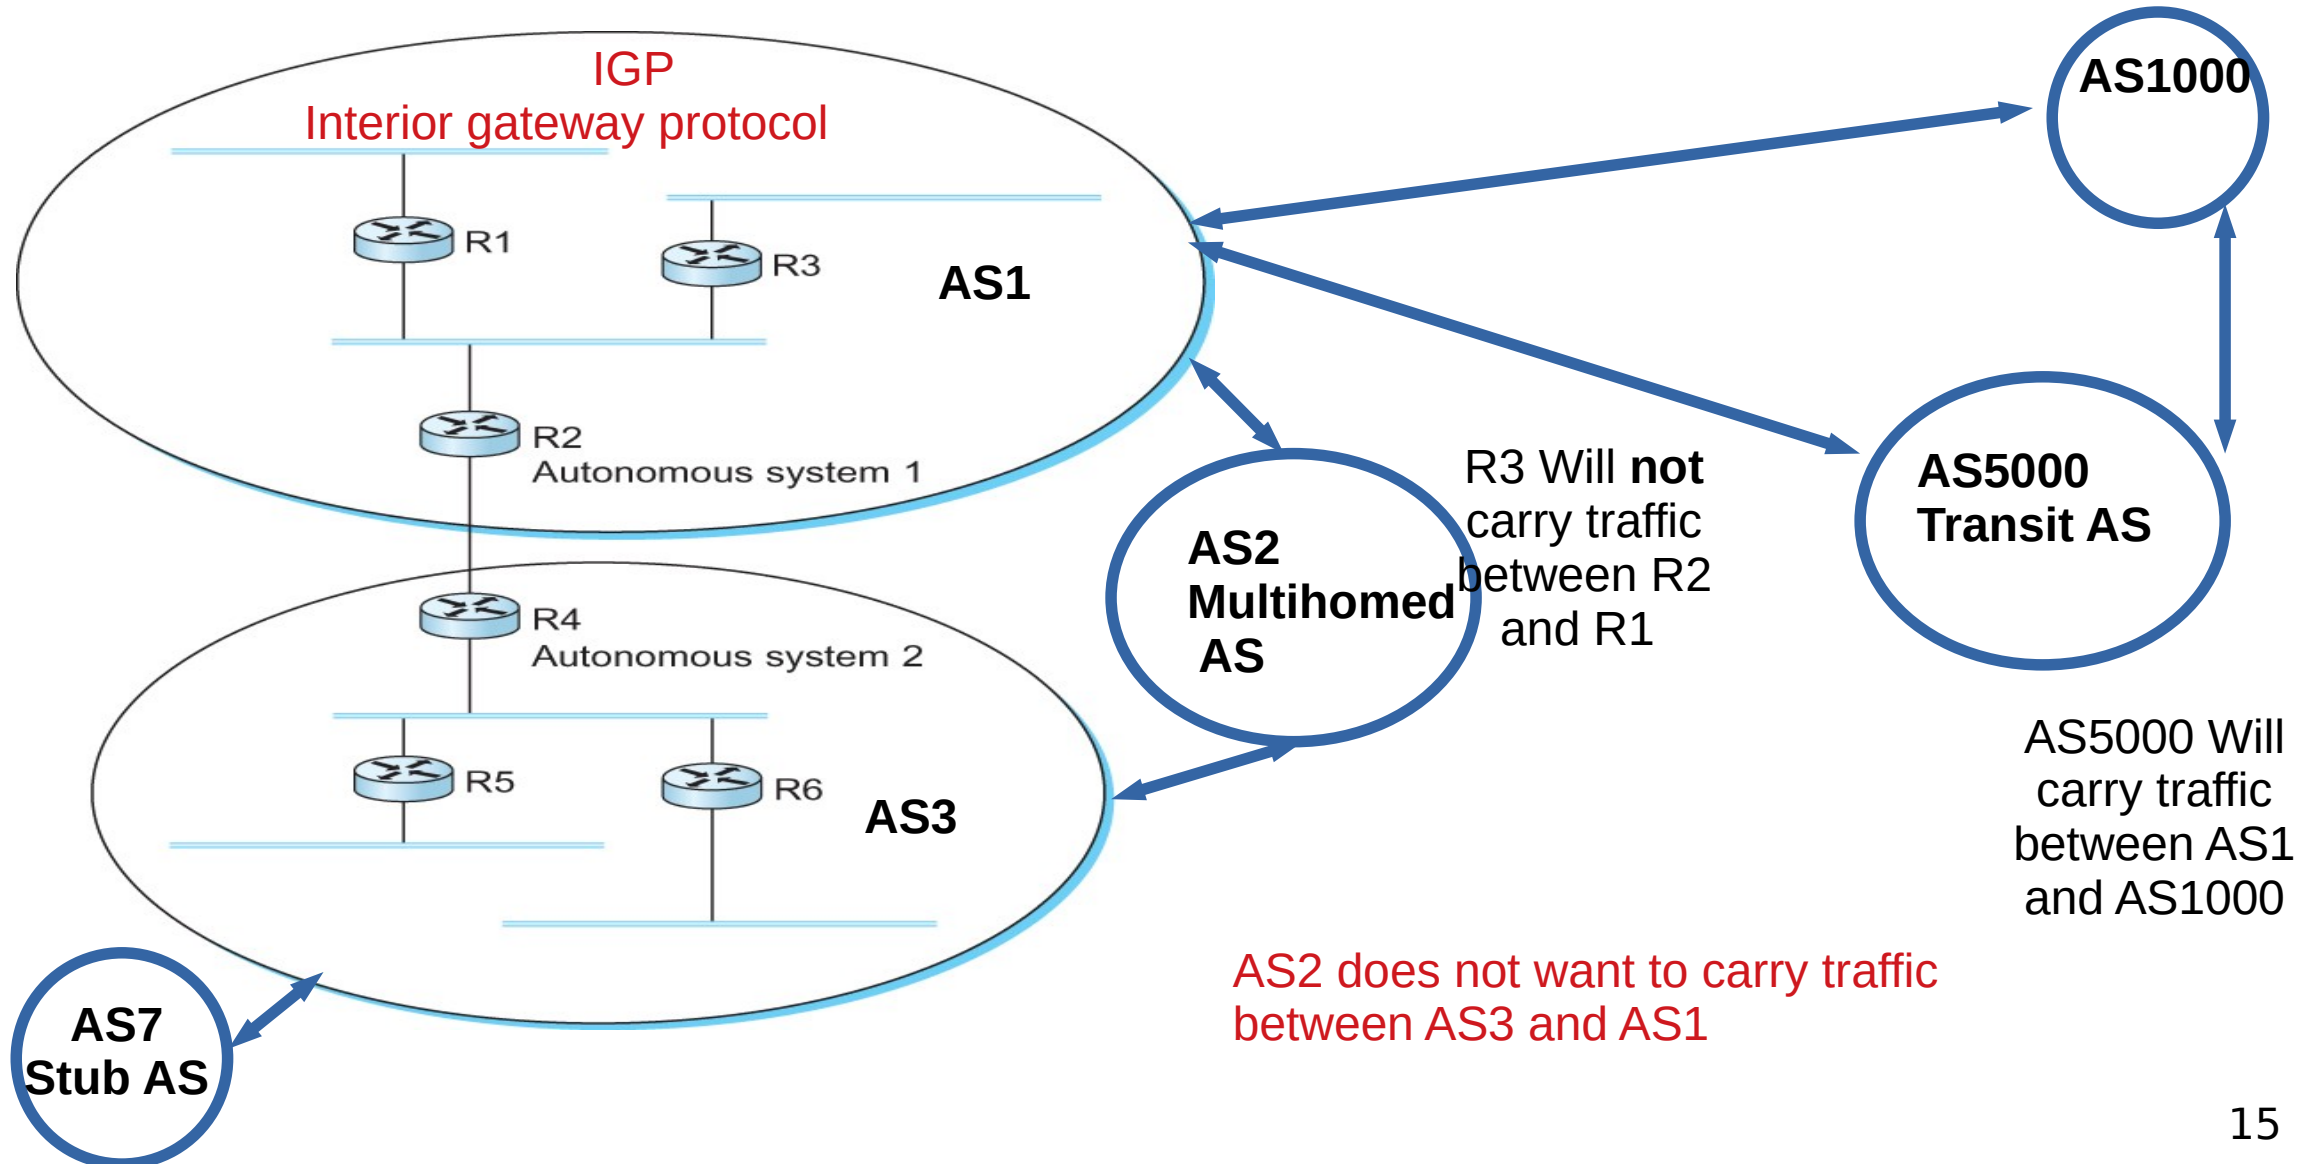
BGP: Allows for policy

- Capable of enforcing various policies
 - AS2 → Don't use AS1 to get to AS3
- Not part of BGP – configuration information that controls propagation of paths

BGP: Hop by Hop model and control what you tell your neighbors

- You can only tell others what you are using
 - But you control what you say
- BGP advertises only to peers
 - Tell them what you are using
 - Hop-by-hop model

What should AS2 (multihomed) tell AS3?



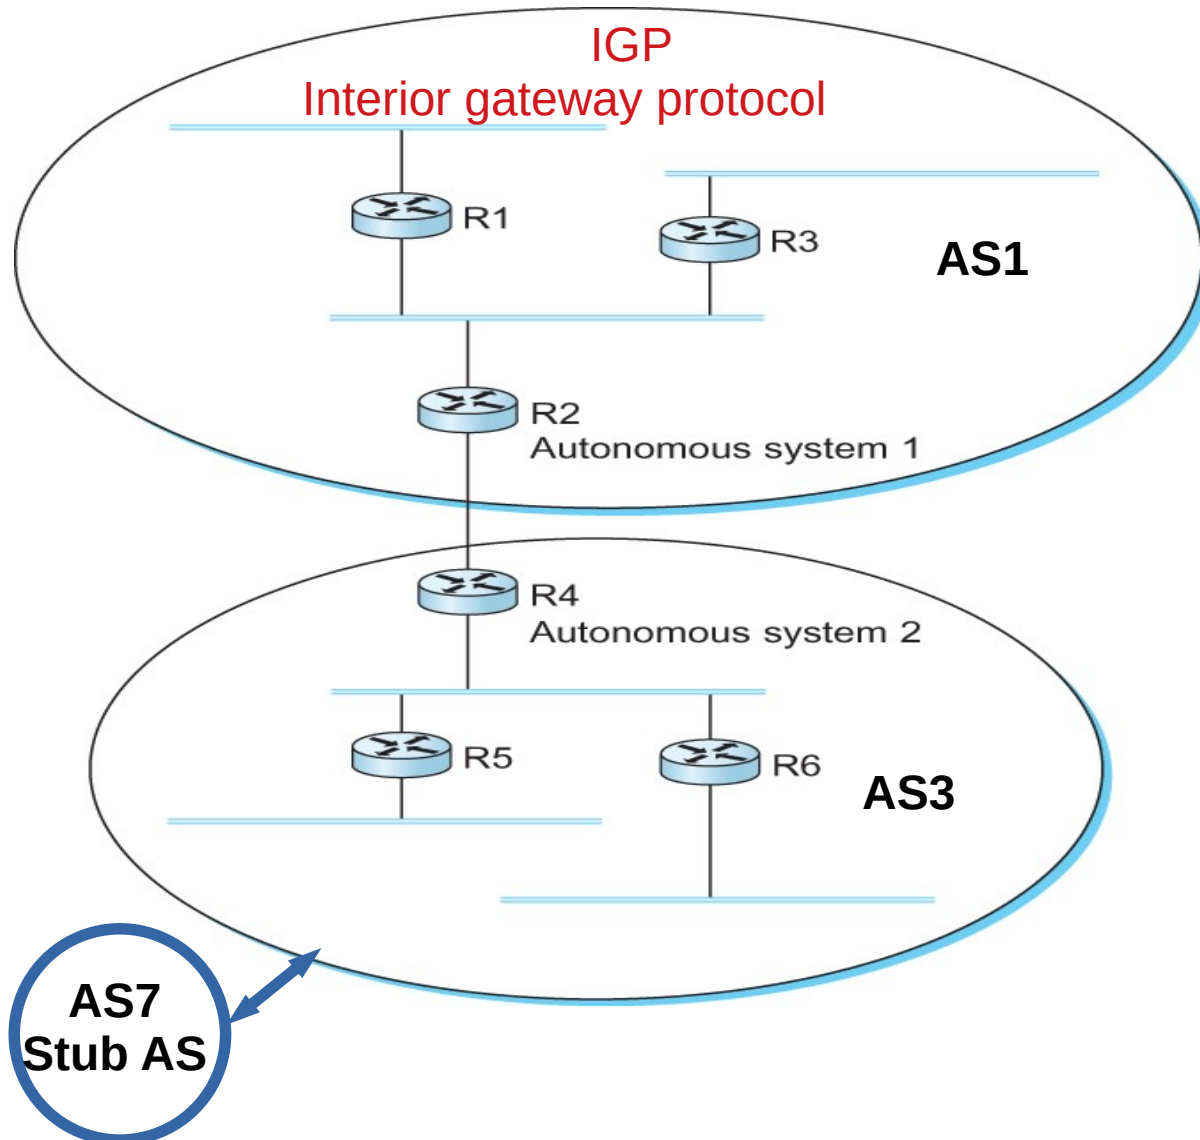
Examples BGP Policies

- Multihomed AS100 does not want to act as a transit
 - Limit advertisement
- If someone pays AS100 – yes
 - Advertise only to those who are paying
- Prefer one path over the other
 - Play with the cost, artificially increase path length and so on ← more on this later

Examples BGP Policies

- Multihomed AS100 does not want to act as a transit
 - Limit advertisement
- If someone pays AS100 – yes
 - Advertise only to those who are paying
- Prefer one path over the other
 - Play with the cost, artificially increase path length and so on ← more on this later

You don't need BGP for Stub ASes

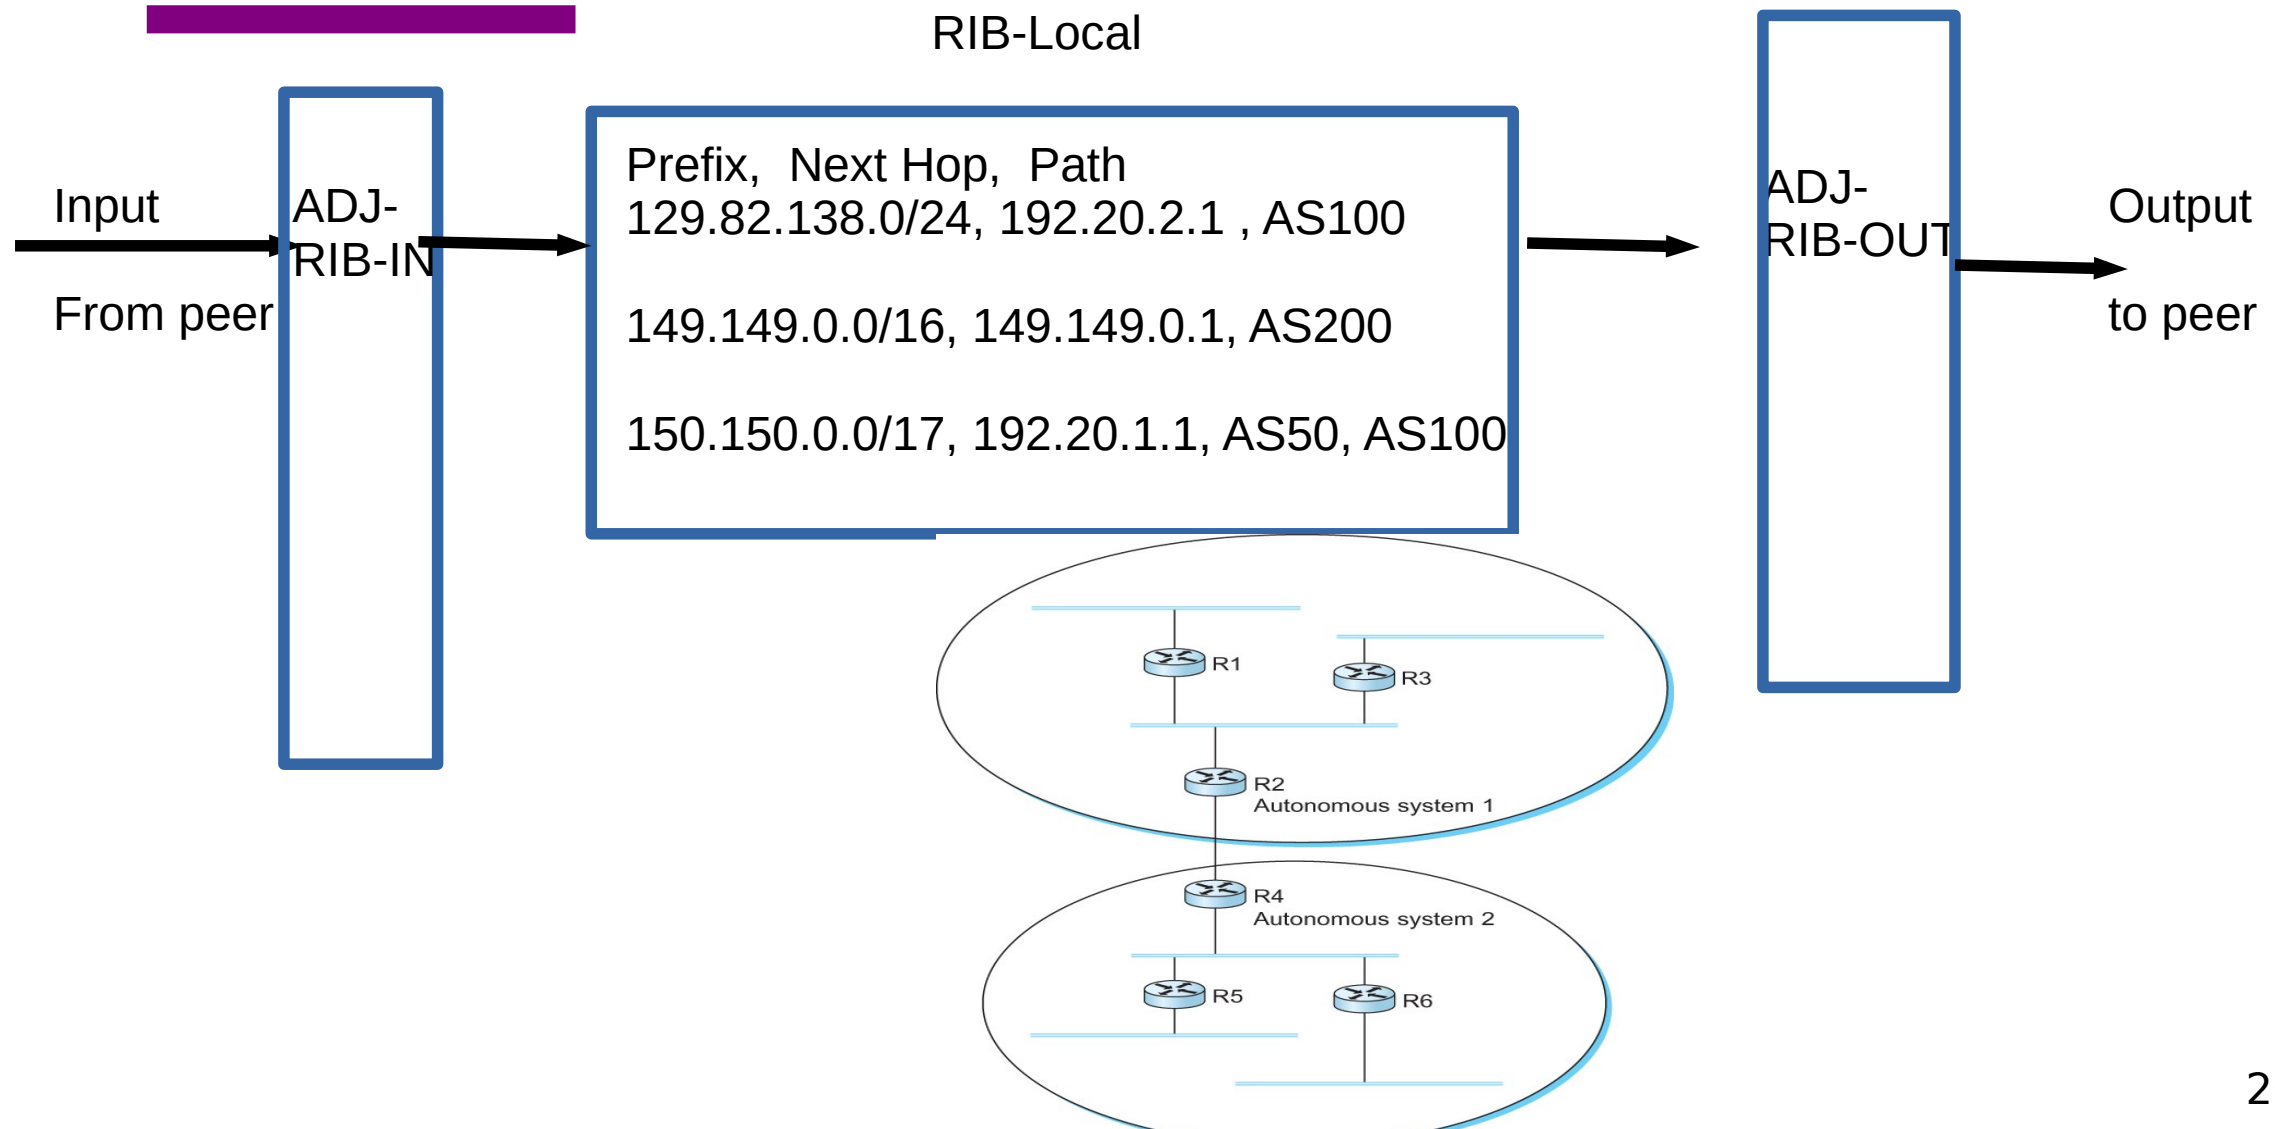


Default IP route should be sufficient

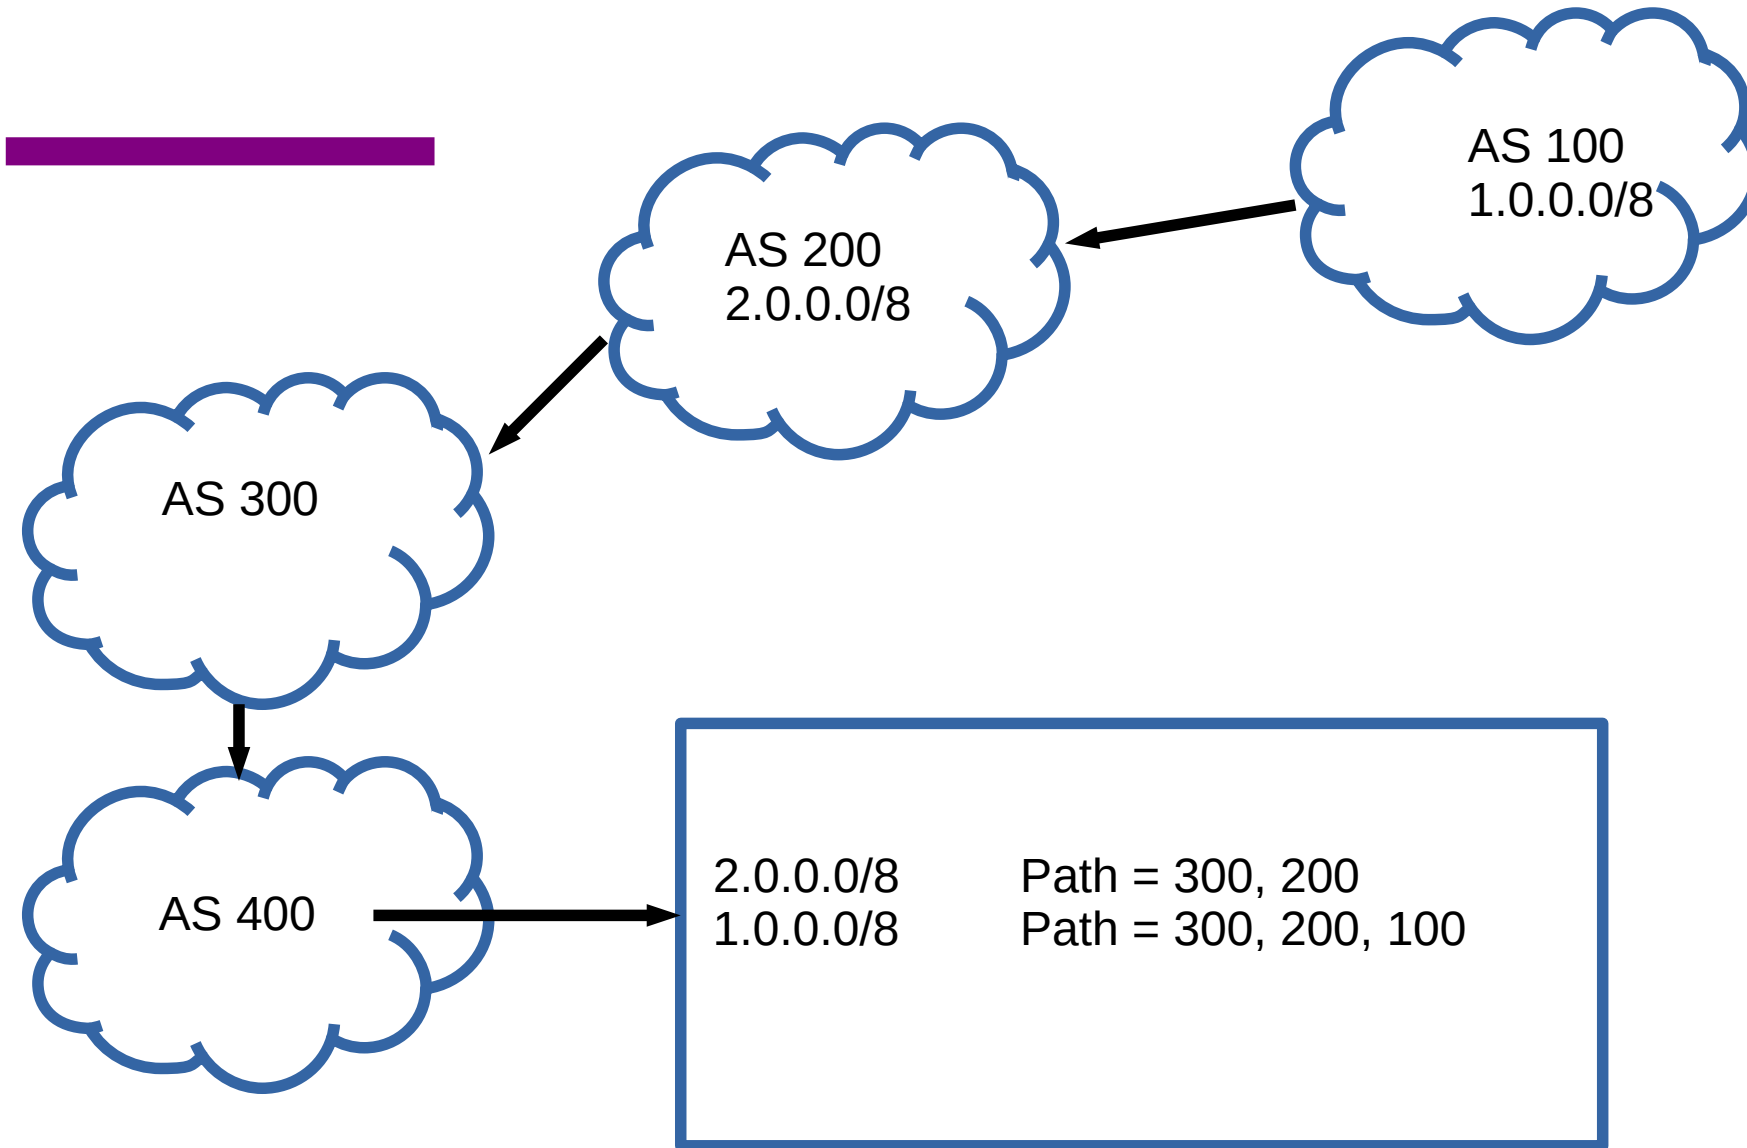
BGP Messages

- Open – Open a TCP connection to a peer
- Update – Update route attributes or withdraw
- Notification – Error notification, close connection
- Keep alive – Periodic update to peers

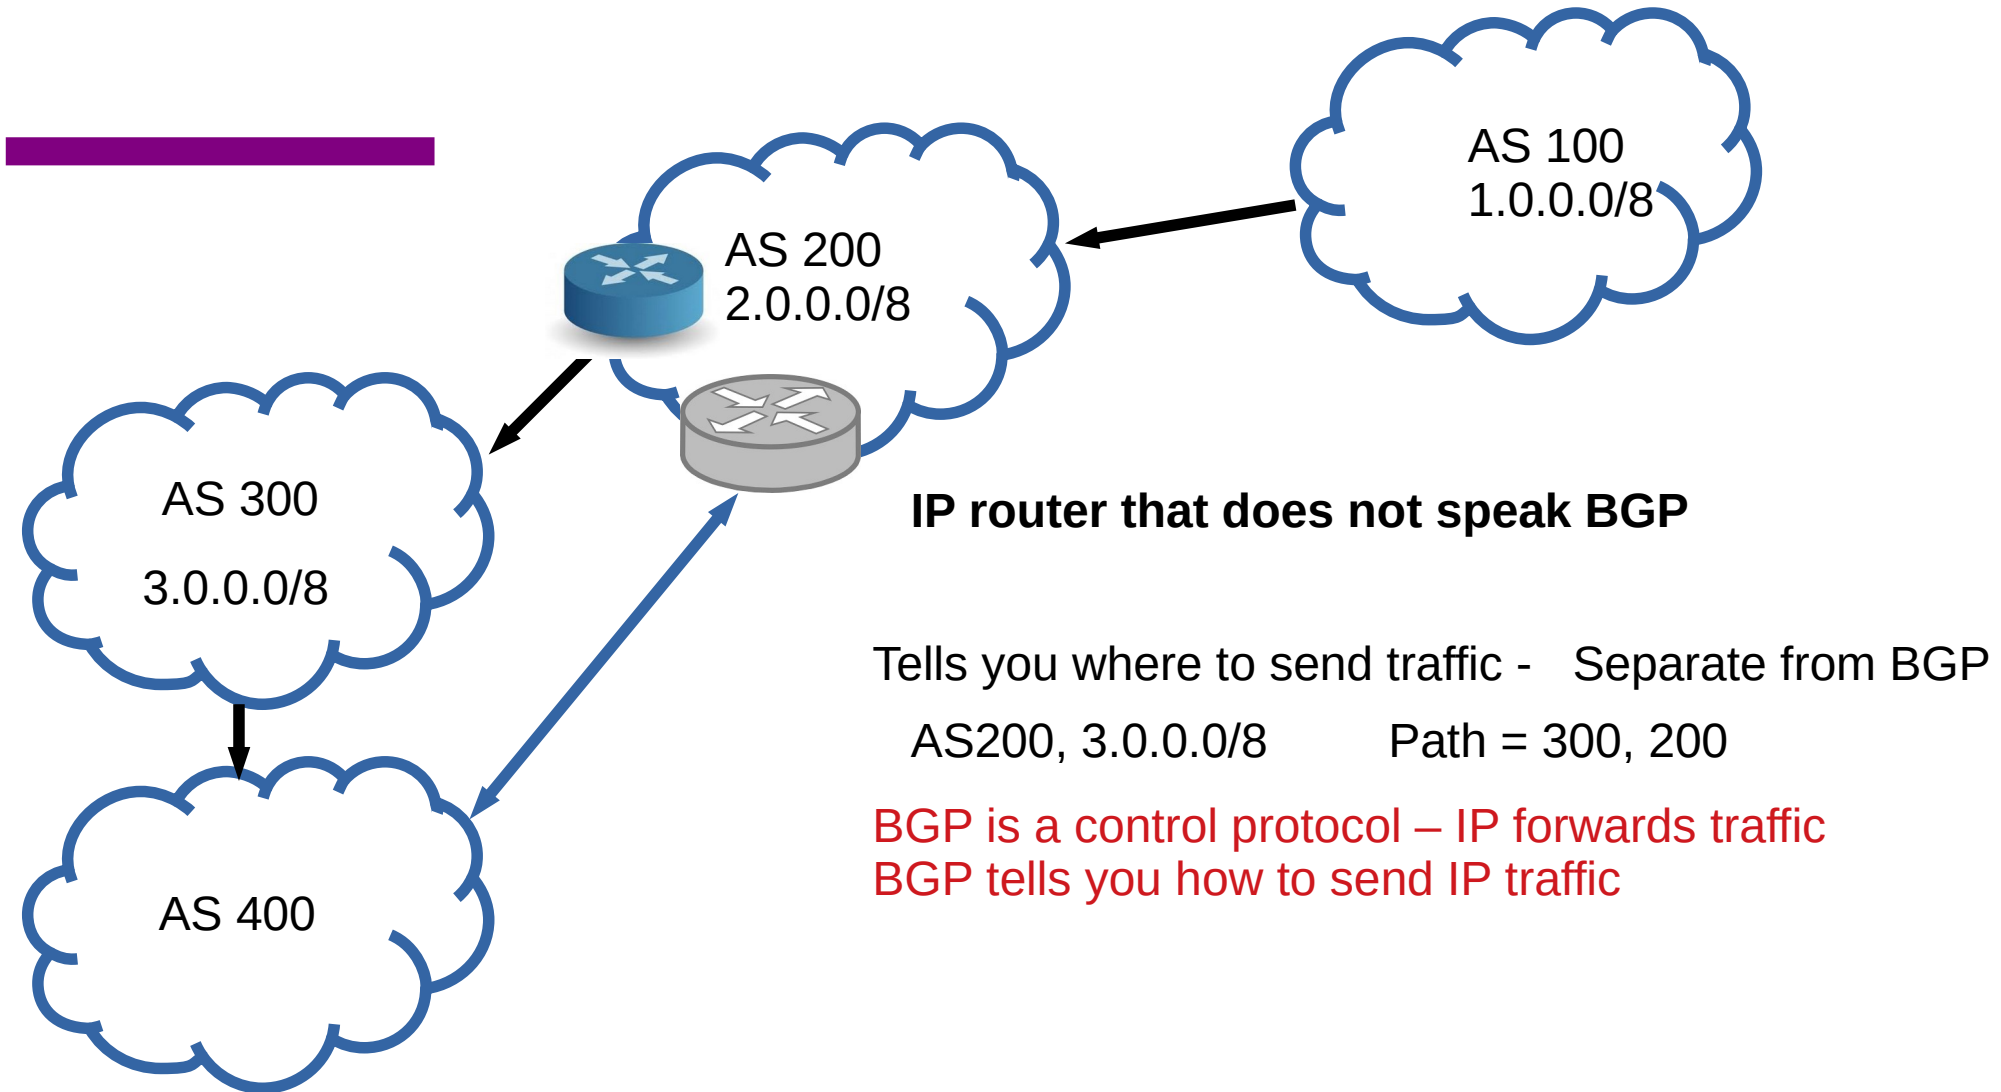
Routing Information Bases (RIB)



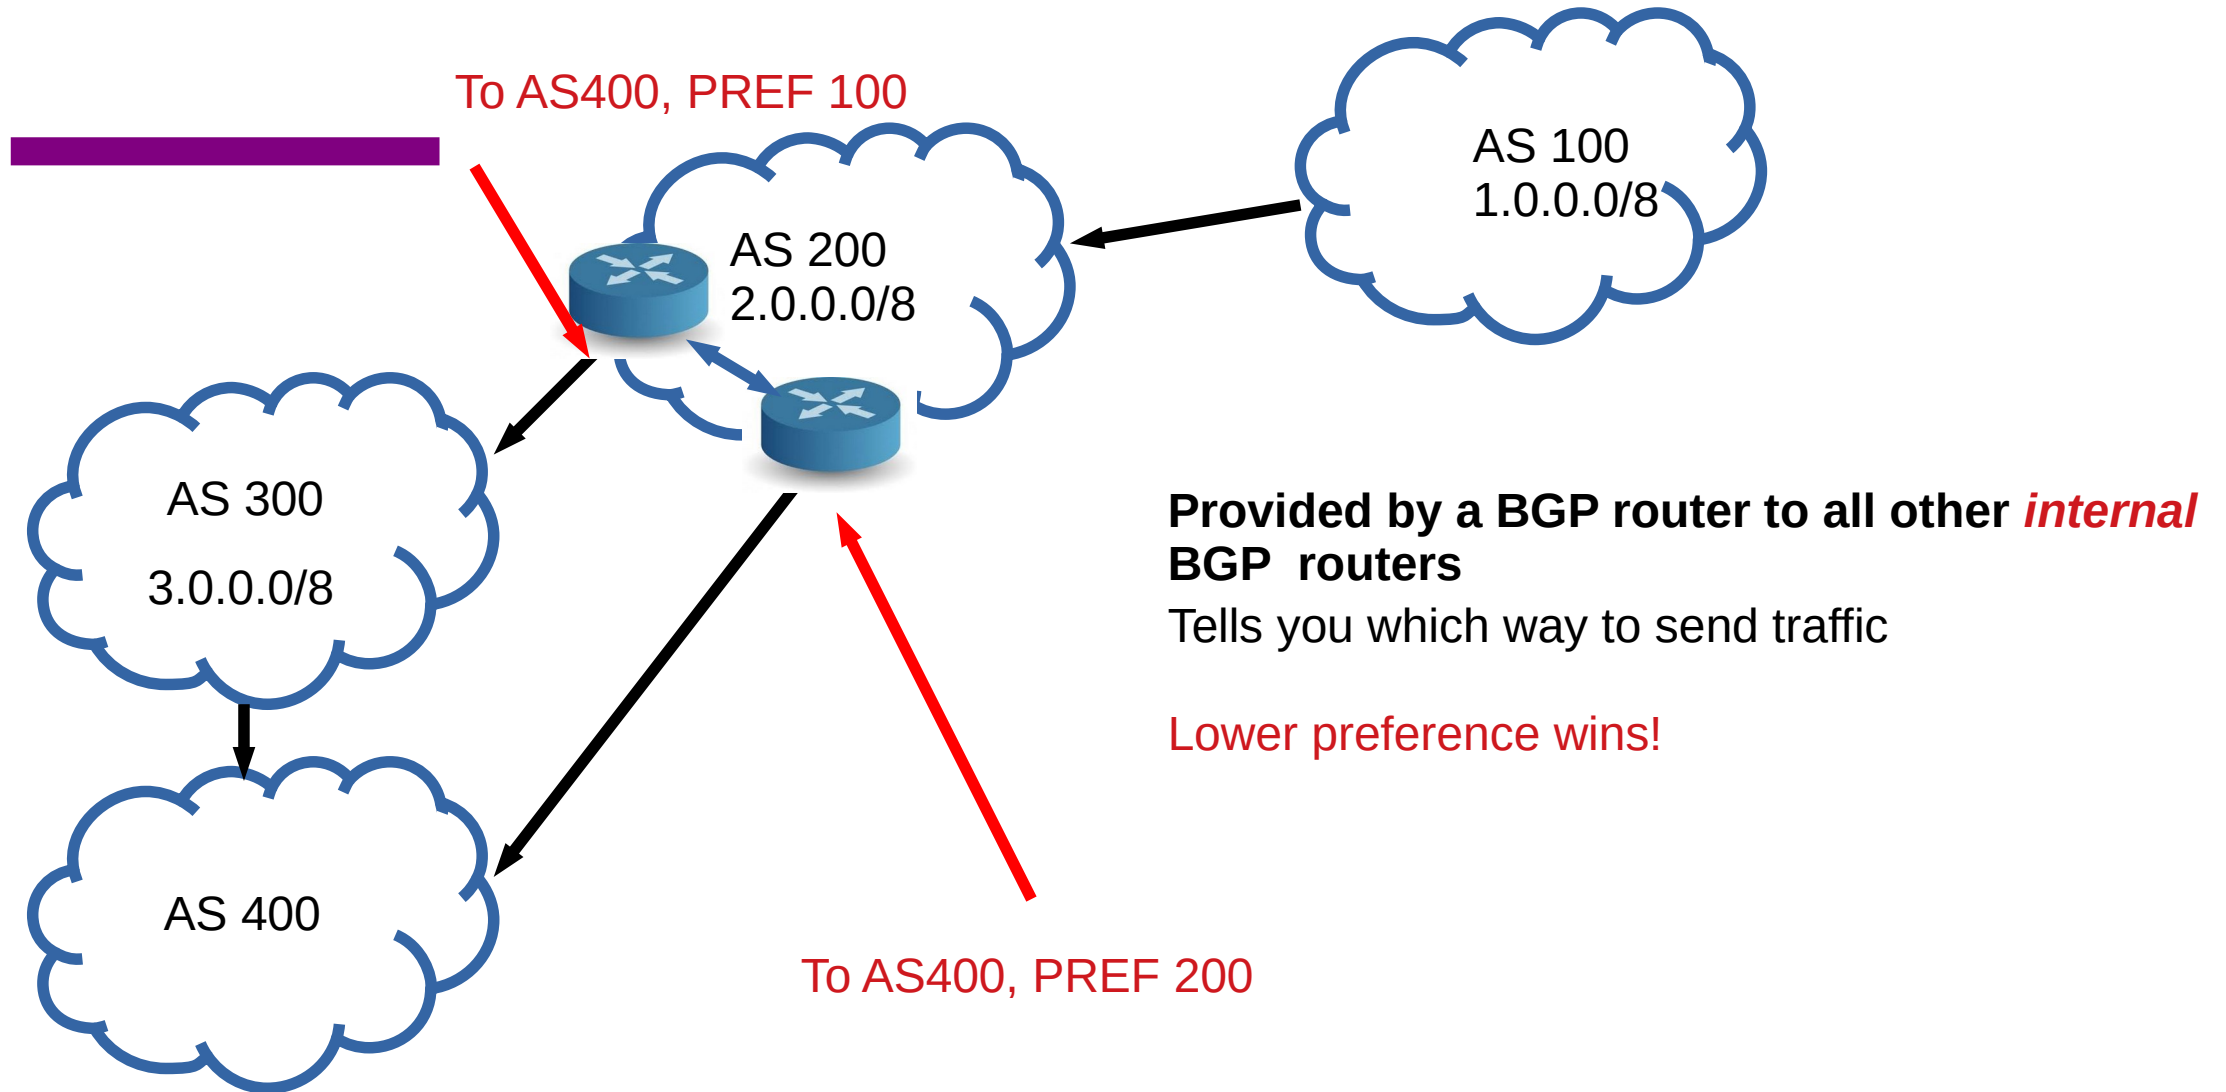
BGP Attributes - AS_PATH




BGP Attributes – Next hop?



BGP Attributes – LOCAL-PREF



BGP Attributes – LOCAL-PREF Example



How to force all traffic to 400 over path b:d?

Set local preference 100 at b
Set local preference 200 at a